

Here's a Python script outlining the implementation of the Mode-Conditional Bayesian MAML framework, incorporating key strategies for efficiency and speed as described in the paper:

Python

```
import torch
import torch.nn as nn
import torch.optim as optim
from torch.distributions import Normal
import numpy as np

# 1. Mode Identification Module (Task Encoder) [cite: 3, 19]
class TaskEncoder(nn.Module):
    def __init__(self, input_dim, embedding_dim):
        super(TaskEncoder, self).__init__()
        # A simple lightweight architecture [cite: 35]
        self.encoder = nn.Sequential(
            nn.Linear(input_dim, 64),
            nn.ReLU(),
            nn.Linear(64, embedding_dim)
        )

    def forward(self, support_set):
        # Process the support set to generate a task embedding [cite: 3,
19]
        embedding = self.encoder(support_set)
        return embedding

# 2. Hypernetwork for Prior Generation [cite: 23, 24]
class PriorHyperNetwork(nn.Module):
    def __init__(self, embedding_dim, prior_params_dim):
        super(PriorHyperNetwork, self).__init__()
        # Generates parameters (mean and log variance) of the prior
distribution [cite: 23, 24]
        self.hypernetwork = nn.Sequential(
            nn.Linear(embedding_dim, 128),
            nn.ReLU(),
            nn.Linear(128, prior_params_dim)
        )

    def forward(self, task_embedding):
        # Output the parameters of the prior distribution [cite: 23, 24]
        prior_params = self.hypernetwork(task_embedding)
        return prior_params
```

```

# 3. Mode-Conditional Bayesian MAML Framework [cite: 76, 77, 78, 79]
class ModeConditionalBayesianMAML(nn.Module):
    def __init__(self, task_encoder, prior_hypernetwork, model):
        super(ModeConditionalBayesianMAML, self).__init__()
        self.task_encoder = task_encoder
        self.prior_hypernetwork = prior_hypernetwork
        self.model = model # The base model

    def meta_train(self, meta_train_data, optimizer, meta_objective):
        """
        Meta-training loop for Mode-Conditional Bayesian MAML. [cite: 25,
        26, 27, 28, 29, 30]

        Args:
            meta_train_data: A list of tasks, where each task is a tuple
            of (support_set, query_set).
            optimizer: The optimizer for meta-training.
            meta_objective: The meta-objective function (loss)
        """
        for task in meta_train_data:
            support_set, query_set = task

            # 1. Mode Identification [cite: 30, 31, 32]
            task_embedding = self.task_encoder(support_set)

            # 2. Prior Distribution Generation [cite: 31, 32]
            prior_params = self.prior_hypernetwork(task_embedding)
            # Assuming prior is a Gaussian distribution [cite: 43]
            prior_mean = prior_params[:self.get_model_params_dim()]
            prior_log_var = prior_params[self.get_model_params_dim():]
            prior_std = torch.exp(0.5 * prior_log_var)
            prior_dist = Normal(prior_mean, prior_std)

            # 3. Bayesian Adaptation [cite: 32, 33, 34]
            # Sample parameters from the prior [cite: 32, 33, 34]
            params = prior_dist.rsample() # Sample parameters
            self.set_model_params(params) # Set the sampled params to the
            model

            # Perform adaptation (e.g., gradient descent or Stein
            Variational Gradient Descent (SVGD)) [cite: 37, 38, 39, 40, 41]
            # For simplicity, using gradient descent here
            adapted_params = self.adapt_model(support_set, self.model)

            # 4. Evaluate on Query Set [cite: 49, 50]

```

```

        query_loss = meta_objective(self.model, adapted_params,
query_set) # calculate loss with the adapted params

        # 5. Meta-Update [cite: 49, 50]
        optimizer.zero_grad()
        query_loss.backward()
        optimizer.step()

def meta_test(self, test_task):
    """
    Meta-testing loop for Mode-Conditional Bayesian MAML. [cite: 30,
31, 32, 33, 34]

    Args:
        test_task: A tuple of (support_set, query_set) for the test
task.
    """
    support_set, query_set = test_task

    # 1. Mode Identification [cite: 30, 31, 32]
    task_embedding = self.task_encoder(support_set)

    # 2. Prior Distribution Generation [cite: 31, 32]
    prior_params = self.prior_hypernetwork(task_embedding)
    # Assuming prior is a Gaussian distribution [cite: 43]
    prior_mean = prior_params[:self.get_model_params_dim()]
    prior_log_var = prior_params[self.get_model_params_dim():]
    prior_std = torch.exp(0.5 * prior_log_var)
    prior_dist = Normal(prior_mean, prior_std)

    # 3. Bayesian Inference / Adaptation [cite: 32, 33, 34]
    # Sample parameters from the prior [cite: 32, 33, 34]
    params = prior_dist.rsample()
    self.set_model_params(params)

    # Perform adaptation (e.g., gradient descent) [cite: 37, 38, 39,
40, 41]
    adapted_params = self.adapt_model(support_set, self.model)

    # 4. Make Predictions [cite: 32, 33, 34]
    # Sample from posterior or use the adapted parameters to make
predictions [cite: 32, 33, 34]
    # For simplicity, using the adapted parameters directly
    predictions = self.model.forward(query_set[0]) # Assuming
query_set is a tuple (data, labels)
    return predictions

```

```

def adapt_model(self, support_set, model):
    """
    Adapt the model parameters using the support set (e.g., via
    gradient descent). [cite: 37, 38, 39, 40, 41]

    Args:
        support_set: The support set for the task.
        model: the model to adapt

    Returns:
        Adapted parameters.
    """
    # Simple gradient descent adaptation [cite: 37, 38, 39, 40, 41]

    support_data, support_labels = support_set

    optimizer = optim.Adam(model.parameters(), lr=0.01) # Example
optimizer

    for _ in range(5): # Adaptation steps
        optimizer.zero_grad()
        outputs = model(support_data)
        loss = nn.functional.cross_entropy(outputs, support_labels)
        loss.backward()
        optimizer.step()

    adapted_params = self.get_model_params()
    return adapted_params

def get_model_params(self):
    """
    Get all parameters of the model as a single vector.
    """
    return torch.cat([param.view(-1) for param in
self.model.parameters()])

def set_model_params(self, params):
    """
    Set the model parameters using a single vector.
    """
    params_ptr = 0
    for param in self.model.parameters():
        num_param = param.numel()

```

```

        param.data.copy_(params[params_ptr:params_ptr +
num_param].view_as(param))
        params_ptr += num_param

def get_model_params_dim(self):
    """
    Calculates the total number of parameters in the model.
    """
    return sum(param.numel() for param in self.model.parameters())

# Example Usage [cite: 68, 69, 70, 71, 72, 73, 74]

# 1. Define the base model
class BaseModel(nn.Module):
    def __init__(self, input_dim, output_dim):
        super(BaseModel, self).__init__()
        self.layer = nn.Linear(input_dim, output_dim)

    def forward(self, x):
        return self.layer(x)

input_dim = 10
output_dim = 5
model = BaseModel(input_dim, output_dim)

# 2. Define Task Encoder and Prior Hypernetwork
embedding_dim = 8
prior_params_dim = model.layer.weight.numel() + model.layer.bias.numel() #
Mean and log variance for each parameter [cite: 23, 24]
task_encoder = TaskEncoder(input_dim, embedding_dim)
prior_hypernetwork = PriorHyperNetwork(embedding_dim, prior_params_dim *
2) # times 2 for mean and log variance [cite: 23, 24]

# 3. Initialize Mode-Conditional Bayesian MAML
mc_bamaml = ModeConditionalBayesianMAML(task_encoder, prior_hypernetwork,
model)

# 4. Define Optimizer and Meta-Objective Function
optimizer = optim.Adam(mc_bamaml.parameters(), lr=0.001)

def meta_objective(model, adapted_params, query_set):
    """
    Example meta-objective function (mean squared error). [cite: 49, 50]
    """
    query_data, query_labels = query_set
    predictions = model(query_data)
    loss = nn.functional.cross_entropy(predictions, query_labels)

```

```

        return loss

# 5. Generate Dummy Data for Meta-Training
meta_train_data =
for _ in range(100):
    support_data = torch.randn(5, input_dim)
    support_labels = torch.randint(0, output_dim, (5,))
    query_data = torch.randn(5, input_dim)
    query_labels = torch.randint(0, output_dim, (5,))
    meta_train_data.append(((support_data, support_labels), (query_data,
query_labels)))

# 6. Perform Meta-Training
mc_bamaml.meta_train(meta_train_data, optimizer, meta_objective)

# 7. Generate Dummy Data for Meta-Testing
support_data = torch.randn(5, input_dim)
support_labels = torch.randint(0, output_dim, (5,))
query_data = torch.randn(5, input_dim)
query_labels = torch.randint(0, output_dim, (5,))
meta_test_task = ((support_data, support_labels), (query_data,
query_labels))

# 8. Perform Meta-Testing
predictions = mc_bamaml.meta_test(meta_test_task)
print("Predictions:", predictions)

```

Explanation and Key Points:

- Mode Identification (Task Encoder):** The `TaskEncoder` class is a simple neural network that takes a support set as input and generates a task embedding. This embedding represents the task's mode.
- Prior Hypernetwork:** The `PriorHyperNetwork` class takes the task embedding and outputs the parameters of the prior distribution (mean and log variance for a Gaussian).
- Mode-Conditional Bayesian MAML Class:**

- `meta_train()`: Implements the meta-training loop. It samples tasks, identifies their modes, generates mode-conditional priors, performs Bayesian adaptation (using gradient descent for simplicity), and updates the model.
- `meta_test()`: Implements the meta-testing loop. It identifies the mode of a new task, generates the corresponding prior, adapts to the task, and makes predictions.
- `adapt_model()`: Performs model adaptation using gradient descent on the support set.

4. Efficiency Considerations:

- Lightweight Task Encoder: A simple `nn.Sequential` model is used for the task encoder.
- Hypernetworks for Prior Generation: Hypernetworks are used to generate prior parameters, which is efficient for continuous task spaces.
- Simplified Posterior Approximation: The code assumes a Gaussian posterior distribution for simplicity.

5. Bayesian Adaptation:

- The `adapt_model` function now uses a basic gradient descent approach for adaptation. In a full implementation, you might replace this with a more sophisticated Bayesian inference technique like SVGD.

6. Meta-Objective Function:

- A placeholder `meta_objective` function is defined. In practice, you would use a loss function appropriate to your task.

7. Parameter Handling:

- The `get_model_params`, `set_model_params`, and `get_model_params_dim` methods are used to get and set the model's parameters as a single vector, which is useful for sampling and updating.

8. Example Usage:

- The example code demonstrates how to define a base model, task encoder, and prior hypernetwork, and how to use the `ModeConditionalBayesianMAML` class for meta-training and meta-testing.

Further Improvements:

- **Stein Variational Gradient Descent (SVGD):** Implement SVGD for more efficient and accurate Bayesian inference.
- **Amortized Inference:** Add an inference network to directly predict posterior parameters, bypassing iterative inference during meta-testing.
- **Task Mode Optimization:** Implement strategies for optimizing the definition of task modes, such as task clustering.
- **Parallelization:** Utilize parallel computing techniques to speed up meta-training and meta-testing.
- **More Complex Architectures:** Experiment with more complex architectures for the task encoder, hypernetworks, and base model.
- **Regularization:** Add regularization techniques to prevent overfitting and improve generalization.
- **Evaluation:** Evaluate the framework on standard meta-learning benchmarks.